

```
In [2]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.svm import SVC
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, r2_score

data = pd.read_csv('heart.csv')
print(data)
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	\
0	40	M	ATA	140	289	0	Normal	
1	49	F	NAP	160	180	0	Normal	
2	37	M	ATA	130	283	0	ST	
3	48	F	ASY	138	214	0	Normal	
4	54	M	NAP	150	195	0	Normal	
..	
913	45	M	TA	110	264	0	Normal	
914	68	M	ASY	144	193	1	Normal	
915	57	M	ASY	130	131	0	Normal	
916	57	F	ATA	130	236	0	LVH	
917	38	M	NAP	138	175	0	Normal	

	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	172	N	0.0	Up	0
1	156	N	1.0	Flat	1
2	98	N	0.0	Up	0
3	108	Y	1.5	Flat	1
4	122	N	0.0	Up	0
..
913	132	N	1.2	Flat	1
914	141	N	3.4	Flat	1
915	115	Y	1.2	Flat	1
916	174	N	0.0	Flat	1
917	173	N	0.0	Up	0

[918 rows x 12 columns]

```
In [3]: numerical_cols = ['RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak']
for col in numerical_cols:
    data[col] = data[col].replace(0, np.nan)
    data[col] = data[col].fillna(data[col].median())
```

```
In [4]: categorical_cols = ['Sex', 'ChestPainType', 'RestingECG', 'ExerciseAngina', 'ST_Slo
label_encoders = {}
for col in categorical_cols:
    le = LabelEncoder()
    data[col] = le.fit_transform(data[col])
    label_encoders[col] = le
```

```
In [5]: X = data.drop('HeartDisease', axis=1)
        y = data['HeartDisease']
```

```
In [6]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

        # Scale the features
```

```
In [7]: scaler = StandardScaler()
        X_train = scaler.fit_transform(X_train)
        X_test = scaler.transform(X_test)
```

```
In [8]: svm_model = SVC(random_state=42)
        lr_model = LogisticRegression(random_state=42) # Logistic Regression for classification
        linear_reg = LinearRegression() # Linear Regression for experimental comparison
        dt_model = DecisionTreeClassifier(random_state=42)
        rf_model = RandomForestClassifier(random_state=42)
```

```
In [9]: svm_model.fit(X_train, y_train)
        lr_model.fit(X_train, y_train)
        linear_reg.fit(X_train, y_train) # Linear Regression
        dt_model.fit(X_train, y_train)
        rf_model.fit(X_train, y_train)
```

```
Out[9]: ▼      RandomForestClassifier
        RandomForestClassifier(random_state=42)
```

```
In [10]: svm_pred = svm_model.predict(X_test)
          lr_pred = lr_model.predict(X_test)
          linear_reg_pred = linear_reg.predict(X_test) # Continuous output
          dt_pred = dt_model.predict(X_test)
          rf_pred = rf_model.predict(X_test)
```

```
In [11]: linear_reg_pred_binary = (linear_reg_pred >= 0.5).astype(int)
```

```
In [12]: print("Support Vector Machine Accuracy:", accuracy_score(y_test, svm_pred))
          print("\nSVM Classification Report:\n", classification_report(y_test, svm_pred))
```

Support Vector Machine Accuracy: 0.8369565217391305

SVM Classification Report:

	precision	recall	f1-score	support
0	0.79	0.83	0.81	77
1	0.87	0.84	0.86	107
accuracy			0.84	184
macro avg	0.83	0.84	0.83	184
weighted avg	0.84	0.84	0.84	184

```
In [13]: print("Logistic Regression Accuracy:", accuracy_score(y_test, lr_pred))
```

```
print("Logistic Regression Classification Report:\n", classification_report(y_test,
```

Logistic Regression Accuracy: 0.842391304347826

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.78	0.87	0.82	77
1	0.90	0.82	0.86	107
accuracy			0.84	184
macro avg	0.84	0.85	0.84	184
weighted avg	0.85	0.84	0.84	184

```
In [14]: print("Linear Regression Accuracy (thresholded):", accuracy_score(y_test, linear_re
print("Linear Regression Classification Report (thresholded):\n", classification_re
print("\nLinear Regression R-squared:", r2_score(y_test, linear_reg_pred))
```

Linear Regression Accuracy (thresholded): 0.8260869565217391

Linear Regression Classification Report (thresholded):

	precision	recall	f1-score	support
0	0.75	0.87	0.81	77
1	0.89	0.79	0.84	107
accuracy			0.83	184
macro avg	0.82	0.83	0.82	184
weighted avg	0.84	0.83	0.83	184

Linear Regression R-squared: 0.4404783267915111

```
In [15]: print("Decision Tree Accuracy:", accuracy_score(y_test, dt_pred))
print("Decision Tree Classification Report:\n", classification_report(y_test, dt_pr
```

Decision Tree Accuracy: 0.7880434782608695

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	0.71	0.84	0.77	77
1	0.87	0.75	0.80	107
accuracy			0.79	184
macro avg	0.79	0.80	0.79	184
weighted avg	0.80	0.79	0.79	184

```
In [16]: print("Random Forest Accuracy:", accuracy_score(y_test, rf_pred))
print("Random Forest Classification Report:\n", classification_report(y_test, rf_pr
```

Random Forest Accuracy: 0.8586956521739131

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.80	0.88	0.84	77
1	0.91	0.84	0.87	107
accuracy			0.86	184
macro avg	0.85	0.86	0.86	184
weighted avg	0.86	0.86	0.86	184

```
In [17]: import matplotlib.pyplot as plt

models = [
    "Support Vector Machine",
    "Logistic Regression",
    "Linear Regression (Thresholded)",
    "Decision Tree",
    "Random Forest"
]

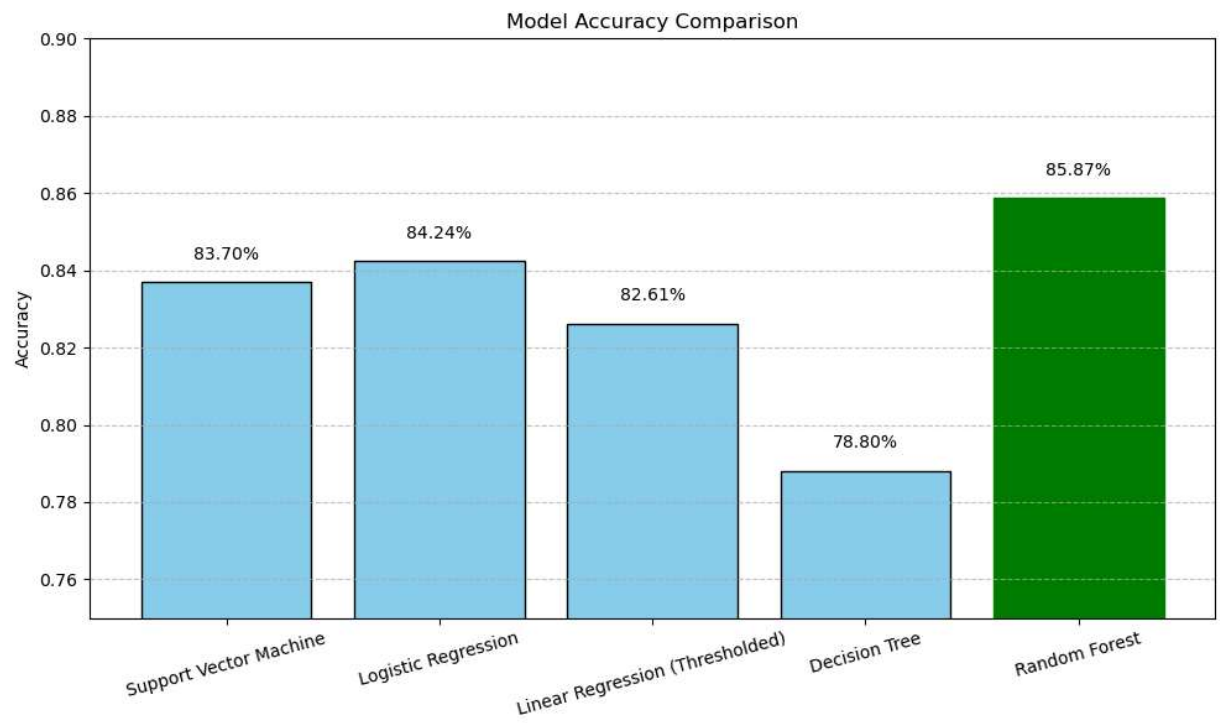
accuracies = [
    0.8369565217391305,
    0.842391304347826,
    0.8260869565217391,
    0.7880434782608695,
    0.8586956521739131
]

plt.figure(figsize=(10, 6))
bars = plt.bar(models, accuracies, color='skyblue', edgecolor='black')
plt.ylim(0.75, 0.9)
plt.ylabel('Accuracy')
plt.title('Model Accuracy Comparison')

best_index = accuracies.index(max(accuracies))
bars[best_index].set_color('green')

for bar in bars:
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width() / 2.0, yval + 0.005, f'{yval:.2%}', ha='center')

plt.xticks(rotation=15)
plt.tight_layout()
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```



In []: